

Judge Panel

Run #47

Four-step multi-rater protocol — coordinate labeling, measure agreement, adjudicate disagreement, escalate to tiebreaker. Kappa floors are enforced before any judge result is allowed to gate.

WORKFLOW · MULTI-RATER PROTOCOL

Currently on **Step 2 / 4** · Step 3 unlocks when every judge is at 60/60

✓ STEP 1 · COMPLETE

Blind independent labeling
All judges labeled 60 of 60 assigned traces in isolation — no cross-talk, no shared session.

2 STEP 2 · IN PROGRESS

Agreement measurement
Cohen's κ across judge pairs — in progress. Waiting on Audio Judge to complete labeling pass.

3 STEP 3 · PENDING

Disagreement review
8 traces queued for adjudication once kappas land. Most are **spec questions**, not labeling errors.

4 STEP 4 · PENDING

Tiebreaker escalations
2 pending · SLA: 5 days · **Oldest open: 3 days**

Judge Roster 5 ROLES				Step 1 · 262 / 300 traces labeled	
Jurisdiction, labeling progress, and current inter-rater agreement per judge					
ROLE	JURISDICTION	PROGRESS		KAPPA	
Content Judge	Factual accuracy, tense fidelity, calendar accuracy	✓	60 / 60	—	
Ingestion Judge	Email thread retention, HTML parsing, calendar completeness	✓	60 / 60	—	
Audio Judge	Voice naturalness, host differentiation, interruption recovery	•	41 / 60	—	
Visual / UX Judge	Transcript-audio sync, readability, visual pacing	○	0 / 60	—	
Tiebreaker (Product Lead)	Final authority · zero-tolerance events	2 escalated SLA: 5 days · oldest open 3 days		—	

Kappa Scores

COHEN'S K

Pairs 1 / 4 measured

Step 2 — awaiting Audio Judge completion · 4 of 4 pairs scheduled

JUDGE PAIR	KAPPA	FLOOR	STATUS
Content ↔ Ingestion · email faithfulness	0.81	≥ 0.70	✓ CLEAR
Content ↔ Audio · email faithfulness	— . —	≥ 0.80	⌚ PENDING
Ingestion ↔ Audio · tense collapse	— . —	≥ 0.70	⌚ PENDING
Audio ↔ Visual · transcript sync	— . —	≥ 0.75	⌚ PENDING

Disagreement Queue

STEP 3 · PREVIEW

Avg wait 2.1 days · oldest 3 days

8 traces queued for adjudication · 2 shown — full queue unlocks when all kappas land

⚠ These are specification questions, not labeling errors. Each disagreement below reflects ambiguity in the failure-mode definition. Must resolve before CI-gating any evaluator built on this kappa pass.

TRACE	CONTENT JUDGE LABEL		AUDIO JUDGE EXPECTED	ACTION
eval-2847	CONTENT JUDGE faithfulness_break · tense collapse	vs.	AUDIO JUDGE EXPECTED natural narrative simplification	← SPEC QUESTION Adjudicate →
eval-2851	CONTENT JUDGE faithfulness_break · gap fill	vs.	AUDIO JUDGE EXPECTED gap_fill · acceptable	← SPEC QUESTION Adjudicate →

Periodic recalibration ensures kappa floors stay valid as labeling guidance evolves

NEXT RECALIBRATION DUE

June 1, 2026 • 21d

20 fresh labeled traces required • monthly cadence

🕒🕒 21 DAYS REMAINING

LAST RECALIBRATION

May 1, 2026 • κ run #43

Content Judge κ 0.79 REV • all pairs above floor

✓✓ ALL JUDGES WITHIN TOLERANCE

Schedule Recalibration →