

FIVE DRIFT MONITOR

Huxe · the five places things can quietly break

↓ Export

⚙ Configure

OVERALL DRIFT POSTURE

2 high2 medium1 low

2

HIGH RISK

Input · Consumption

2

MEDIUM RISK

Ingestion · Model

1

LOW RISK

Judge

D1 / 5

INPUT DRIFT

— variability in real-world inputs

● HIGH RISK

WHAT IT IS

Variability in real-world inputs — user diversity, email formats, thread depth, language patterns.

PRODUCT-SPECIFIC OBSERVATION

Every Huxe user has a completely different email mix. *This is not drift over time — it is baseline variability from day 1. Standard drift detection does not apply.* Profile-stratified coverage (Profile A/B/C) is the correct approach, not drift monitoring.

DETECTION

Stratified profile coverage + user segment breakdown across A · B · C

LAST CHECKED 11 May

NEXT DUE 18 May

MECHANISM Weekly

● ACTIVE

View Traces by Profile →

View Input Distribution →

PROFILE COVERAGE

11 MAY

A

78%

B

54%

C

31%

⚠ Profile C under-represented — 9 of 30 traces. Bump in next sample.

D2 / 5

INGESTION DRIFT

— silent breakage from upstream API changes

● MEDIUM RISK

WHAT IT IS

Silent breakage from upstream API changes — Gmail/Outlook HTML rendering, email API format shifts — without any code change on Huxe's side.

DETECTION MECHANISM

Canary corpus · 20 fixed email documents · ground truth fixed · daily check at 06:00 AM

LAST RUN 11 May · 06:00

SCHEDULE Daily 06:00

OWNER Ingestion · Bot

● ACTIVE

⚠ +3 NEW canary failures vs. yesterday — investigate before next eval run. Most likely: Gmail HTML wrapper update at ~03:00 UTC.

View Canary Details →

Configure Corpus →

Diff vs. Yesterday →

CANARY CORPUS

20 DOCS

Pass / Fail today

14 · 6

14

6

01020304050607080910

11121314151617181920

■ PASS

■ FAIL

■ NEW TODAY

Δ vs. yesterday

+3 NEW

D3 / 5

MODEL DRIFT

— silent provider updates to LLM & TTS

● MEDIUM RISK

WHAT IT IS

LLM and TTS providers silently update their models. Pronunciation and generation behavior change without warning.

DETECTION

Monthly kappa recalibration · 20 fresh labeled traces per judge

SCHEDULE First Monday of each month

LAST 01 May

NEXT 01 Jun

CONTENT JUDGE K 0.79 ✓

● ACTIVE

View Kappa History →

Schedule June Recalibration →

KAPPA TREND

LAST 6 MONTHS

κ = 0.70 floor

DEC

JAN

FEB

MAR

APR

MAY

WHAT IT IS

The LLM powering evaluators is silently updated by its provider. Evaluators become more lenient. Scores look like they're improving when they're not.

FLOOR $\kappa \geq 0.70$

CADENCE Monthly

MECHANISM Auto-demote < floor

● ACTIVE

⚠ Any kappa below floor → demoted to **MONITORING ONLY** (scores logged, not gating).

Recalibrate All →

View Evaluator Audit Log →


EVALUATOR HEALTH			
EVALUATOR	K	FLOOR	STATUS
T-05 Ingestion ingest_judge	0.71	0.70	✓
T-06 Faithfulness not configured	—	—	○
T-07 Factuality not configured	—	—	○
1 active · 2 unconfigured			SIM-EST

WHAT IT IS

Same technical quality, declining engagement. Novelty effect exhausting. Users disengaging even when eval scores are stable.

DETECTION

21-day engagement signal tracking · listen-through & return-next-day



⚠ Listen-through declining. Day 14 → Day 21: -4pp, accelerating. Currently below the 65% target line.

⌚ Minimum 21-day window required before drawing conclusions from this data — at the floor today.

WINDOW 21 days

SAMPLE n = 1,284 sessions

LAST REFRESH 06:00 · 11 May

CONFOUNDERS None controlled

● ACTIVE

View Full Consumption Dashboard →

Register Experiment (T-15) →

Segment by Profile →